

Pause 'n' play: Formalizing asynchronous C[#]

Gavin Bierman

Microsoft Research

Joint work with C. Russo, G. Mainland, E. Meijer and M.
Torgersen

ECOOP, June 2012

What problems do developers face today?



- ▶ Want “fluid” apps
- ▶ i.e. no hourglass/animated circle cursor!

Heart of the problem: **Synchronous operations**

- ▶ Simple to use but:
 - ▶ prevent progress until done,
 - ▶ reveal latency,
 - ▶ waste resources (calling threads).

What did my professor teach me?

Use **asynchronous** operations instead:

- ▶ enable concurrent progress whilst operations are running,
- ▶ hide latency,
- ▶ free up resources (calling threads).

Asynchronous coding has always been possible in C#...
...but it's never been easy.

Today's message: C# 5.0 makes asynchronous programming much easier.

Example: Reading from a stream

Synchronous:

```
int bytesRead =  
    str.Read(...);  
    // read and wait
```

Asynchronous: Task model (2008-)

```
Task<int> task = str.ReadAsync(...); // non-blocking  
// do some work  
int bytesRead = task.Result; // may block
```

Efficient waiting

Efficient waiting uses callbacks, only invoked once done!

```
Task<int> task = str.ReadAsync(...); // non-blocking

task.ContinueWith(doneTask => { // task done!
    int bytesRead = doneTask.Result; // can't block
    ...
});
```

That's the easy bit — the tough part is writing the callback.

Sync vs. async taste challenge



C# 4.0: Sync vs. Async

Synchronous stream length method

```
public static long Length(Stream src) {  
    var buf = new byte[0x1000]; int bytesRead;  
    long totalRead = 0;  
    while ((bytesRead = src.Read(buf, 0, buf.Length)) > 0)  
        totalRead += bytesRead;  
    return totalRead;  
}
```

C# 4.0: Sync vs. Async

Asynchronous stream length method

```
public static Task<long> LengthAsync4(Stream src) {  
    var tcs = new TaskCompletionSource<long>();  
    var buf = new byte[0x1000]; int bytesRead;  
    long totalRead = 0;  
    Action<Task<int>> While = null;  
    While = rdtask => {  
        if ((bytesRead = rdtask.Result) > 0) {  
            totalRead += bytesRead;  
            src.ReadAsync(buf, 0, buf.Length).ContinueWith(While);  
        }  
        else tcs.SetResult(totalRead);  
    };  
    src.ReadAsync(buf, 0, buf.Length).ContinueWith(While);  
    return tcs.Task;  
}
```

"9/10 devs prefer the taste of synchronous code."

The problem with callbacks

Turning a synchronous call into an asynchronous call is hard! (The “Inversion of control” problem.)

Need to capture the next state of the caller as a callback.

Method state includes locals but also implicit control state:

- ▶ program counter (what to do next in this method)
- ▶ runtime stack (caller to return to from this method!)

This is easier in languages with call/cc (Scheme/SMLNJ).

Others have to be clever.

(Honorable mentions: Haskell, F#, Scala)

Sync vs. async taste challenge, again



C# 5.0: Sync vs. Async

Synchronous stream length method

```
public static long Length(Stream src) {  
    var buf = new byte[0x1000]; int bytesRead;  
    long totalRead = 0;  
    while ((bytesRead = src.Read(buf,0,buf.Length))>0)  
        totalRead += bytesRead;  
    return totalRead;}  
}
```

C# 5.0 Async version

```
public static async Task<long> LengthAsync(Stream src) {  
    var buf = new byte[0x1000]; int bytesRead;  
    long totalRead = 0;  
    while ((bytesRead = await src.ReadAsync(buf,0,buf.Length))>0)  
        totalRead += bytesRead;  
    return totalRead;}  
}
```

"Mr C# compiler: please build the callback for me"

"Wow, almost the same—and it's good for me too?"

Async: Statics

```
public static async Task<long> LengthAsync(Stream src) {  
    var buf = new byte[0x1000]; int bytesRead;  
    long totalRead = 0;  
    while ((bytesRead = await src.ReadAsync(buf,0,buf.Length))>0)  
        totalRead += bytesRead;  
    return totalRead;  
}
```

C# 5.0 adds two keywords **async** and **await**:

- ▶ An **async** method must return a **Task<T>** (or **Task** or **void**);
- ▶ ...yet exit by **return** of a **T** (not a **Task<T>**).
- ▶ Only **async** methods can contain await expressions.
- ▶ If **e** has type **Task<U>** then **await e** has type **U**.
- ▶ So one async can await task of another.

Async: Dynamics

```
public static async Task<long> LengthAsync(Stream src) {  
    var buf = new byte[0x1000]; int bytesRead;  
    long totalRead = 0;  
    while ((bytesRead = await src.ReadAsync(buf, 0, buf.Length)) > 0)  
        totalRead += bytesRead;  
    return totalRead;  
}
```

Calling `LengthAsync()`:

- ▶ Creates a fresh **incomplete** task for this call.
- ▶ Executes until the next **await** expression.
- ▶ If awaited argument is **complete** now:
 - ▶ **continue** with result (fast path).
 - ▶ otherwise, **suspend** until completion (slow path).
- ▶ **return** completes this call's task.
- ▶ The first suspend yields the incomplete task to the caller.

Fine points

An async method call:

- ▶ runs synchronously until or unless it needs to suspend.
- ▶ once suspended, may complete asynchronously.
- ▶ if never suspended, completes on same thread.
- ▶ does not (in itself) spawn a new thread.

Synchronous on entry avoids context switching; enables fast path implementations.

Typically, suspensions resume in thread pool or event loop; not on new threads.

Threading details depends on context (it's complicated—we don't formalize it, but we could).

What's the compiler doing?

```
public static Task<long> LengthAsync(Stream src) {  
    var tcs = new TaskCompletionSource<long>(); // tcs.Task new & incomplete  
    var state = 0; TaskAwaiter<int> readAwaiter;  
    byte[] buffer = null; int bytesRead = 0; long totalRead = 0;  
    Action act = null; act = () => {  
        while (true) switch (state++) {  
            case 0: // entry  
                buffer = new byte[0x1000]; totalRead = 0; continue; // goto 1  
            case 1: // while loop at await  
                readAwaiter = src.ReadAsync(buffer, 0, buffer.Length).GetAwaiter();  
                if (readAwaiter.IsCompleted) continue; // continue from 2  
                else { readAwaiter.OnCompleted(act); return; } // suspend at 2  
            case 2: // while loop after await  
                if ((bytesRead = readAwaiter.GetResult()) > 0) {  
                    totalRead += bytesRead;  
                    state = 1; continue; // goto 1  
                }  
                else continue; // goto 3  
            case 3: // while exit  
                tcs.SetResult(totalRead); // complete tcs.Task & "return"  
                return; // exit machine  
        }  
    }; // end of act delegate  
    act(); // start the machine on this thread  
    return tcs.Task;  
} // on first suspend or exit from machine
```

Formal specification

Current C# 5.0 specs:

- ▶ Precise prose describing syntax and typing.
- ▶ Example source to source translations (like previous slide).

Hard to follow, trickier to apply.

Our Goal: a precise, high-level operational semantics for:

- ▶ programmers (perhaps)
- ▶ compiler writers (hopefully)
- ▶ researchers (realistically)

No mention of the finite state machine at all!

But for the non-mathematically inclined...

Summary:

- ▶ We did the proper “featherweight” thing: $FC_5^\#$
- ▶ It works!
 - ▶ **async** can be given a “high-level” operational interpretation
 - ▶ It can be seen as a “simple” extension of $C^\#$ 4.0 semantics
 - ▶ The correctness proof is quite subtle (as we have shared state)

C# 5.0 operational semantics

Three layered small-step relations:

- Frame transitions (frame local steps):

$$H_1 \triangleright F_1 \rightarrow H_2 \triangleright F_2$$

- Frame stack transitions (thread local steps):

$$H_1 \triangleright FS_1 \twoheadrightarrow H_2 \triangleright FS_2$$

- Process transitions (communication):

$$H_1 \triangleright P_1 \leadsto H_2 \triangleright P_2$$

where	F	$::=$	$\langle L, \bar{s} \rangle^\ell$	Frame (locals plus statements)
	L	$::=$	$\{\bar{x} \mapsto \bar{v}\}$	Locals Map (frame state)
	FS	$::=$	$\epsilon \mid F \circ FS$	Frame Stack (thread)
	ℓ	$::=$	s	synchronous frame label
			$a(o)$	asynchronous frame for (task) o
	P	$::=$	$\{FS_1, \dots, FS_n\}$	Process (bag of stacks)

New Stack transitions

Async method call

$$\begin{aligned} H_0 &\triangleright \langle L_0, y_0=y_1.m(\bar{z}); \bar{s} \rangle^\ell \circ FS \\ \rightarrow H_1 &\triangleright \langle L_1, \bar{t} \rangle^{\mathbf{a}(o)} \circ \langle L_0[y_0 \mapsto o], \bar{s} \rangle^\ell \circ FS \\ \text{where } H_0(L_0(y_1)) &= \langle \sigma_1, FM \rangle \\ mbody(\sigma_1, m) &= mb: (\bar{\sigma} \bar{x}) \rightarrow^{\mathbf{a}} \psi, \text{ and } mb = \bar{\tau} \bar{y}; \bar{t} \\ o &\notin dom(H_0) \\ H_1 &= H_0[o \mapsto \langle \psi, \mathbf{running}(\epsilon) \rangle] \\ L_1 &= [\bar{x} \mapsto L_0(\bar{z}), \bar{y} \mapsto default(\bar{\tau}), \mathbf{this} \mapsto L_0(y_1)] \end{aligned}$$

- ▶ An **async** call pushes a new active frame (as before).
- ▶ new frame is labeled asynchronous **a**(*o*).
- ▶ *o* is address of a fresh, running task (no waiters yet).
- ▶ caller will receive task *o* (once active).

New Stack transitions

Async return

$$\begin{aligned} H_0 &\triangleright \{ \langle L, \text{return } y; \bar{s} \rangle^{a(o)} \circ FS \} \cup P \\ &\rightsquigarrow H_1 \triangleright \{ FS \} \cup \text{resume}(\overline{F}) \cup P \\ \text{where } H_0(o) &= \langle \text{Task}\langle \sigma \rangle, \text{running}(\overline{F}) \rangle \\ H_1 &= H_0[o \mapsto \langle \text{Task}\langle \sigma \rangle, \text{done}(L(y)) \rangle] \end{aligned}$$

Returning from an asynchronous method atomically:

- ▶ Completes the frame's task (retrieved from label!)
- ▶ With value "returned".
- ▶ Resumes any suspended frames (as threads).

(Label determines meaning of **return** x)

Await transitions

$$\begin{aligned} H \triangleright \langle L, x=\text{await } y; \bar{s} \rangle^{a(o)} \circ FS \\ \rightarrow H \triangleright \langle L[x \mapsto v], \bar{s} \rangle^{a(o)} \circ FS \\ \text{where } H(L(y)) = \langle \text{Task}\langle \sigma \rangle, \text{done}(v) \rangle \end{aligned}$$

- ▶ **await** on done task continues with its value.

$$\begin{aligned} H_0 \triangleright \langle L, x=\text{await } y; \bar{s} \rangle^{a(o)} \circ FS \rightarrow H_1 \triangleright FS \\ \text{where } L(y) = o_1, H_0(o_1) = \langle \text{Task}\langle \sigma \rangle, \text{running}(\bar{F}) \rangle \\ H_1 = H_0[o_1 \mapsto \langle \text{Task}\langle \sigma \rangle, \text{running}(\langle L, x=y.\text{GetResult}(); \bar{s} \rangle^{a(o)}, \bar{F}) \rangle] \end{aligned}$$

- ▶ **await** on running task adds current frame to its waiters.
- ▶ Waiters call partial **GetResult** which extracts value

Correctness

Usual preservation and progress (modulo usual OO stuck states) technique **but way, WAY, harder**:

- ▶ lots of disjointness requirements on task ids in heap, frame stacks, process
- ▶ tasks satisfy a protocol:
 $\text{running}(\epsilon) \rightarrow^* \text{running}(F, FS) \rightarrow^* \text{done}(v)$

(see paper).

Related work

[Too many to mention!]

Highly relevant:

- ▶ Millstein et al. [TaskJava](#) (PEPM 2007)
 - ▶ similar implementation techniques (state machines).
 - ▶ no pre-emptive concurrency—just convenient, asynchronous sequential execution.
 - ▶ simpler formalization—no heap (!) => no thread communication.
- ▶ F[#]'s asynchronous workflows:
 - ▶ similar ends; different means and tradeoffs.
 - ▶ syntactically heavier (do-notation).
 - ▶ workflows are (inert) values: easier to compose first, run later.
 - ▶ more general (no inherent one-shot restriction).
 - ▶ less efficient (every suspend allocates a new continuation).

[Apologies to other Highly relevant authors☺]

Conclusions

- ▶ C# 5.0 makes it much easier to write asynchronous code:
 - ▶ No need to roll your own callbacks
 - ▶ Builds upon existing **Task** library (which has lots of goodies)
- ▶ The essence of these new features can be captured precisely:
 - ▶ Significant fragment of C# 5.0
 - ▶ Builds upon existing formalization
 - ▶ No mention of finite-state machine encoding
 - ▶ Lots more in the paper:
 1. One-shot semantics
 2. Tail-call optimizations
 3. Awaitable patterns
 4. Exceptions
- ▶ **Future work:** More low-level semantics exposing thread details (including **SynchronizationContext**)

Questions?

